

.dex — Dalvik Executable Format

Copyright © 2007 Google Inc. All rights reserved.

This document describes the layout and contents of .dex files, which are used to hold a set of class definitions and their associated adjunct data.


Guide To Types

| Name | Description |
|-----------|---|
| byte | 8-bit signed int |
| ubyte | 8-bit unsigned int |
| short | 16-bit signed int, little-endian |
| ushort | 16-bit unsigned int, little-endian |
| int | 32-bit signed int, little-endian |
| uint | 32-bit unsigned int, little-endian |
| long | 64-bit signed int, little-endian |
| ulong | 64-bit unsigned int, little-endian |
| sleb128 | signed LEB128, variable-length (see below) |
| uleb128 | unsigned LEB128, variable-length (see below) |
| uleb128p1 | unsigned LEB128 plus 1, variable-length (see below) |

LEB128

LEB128 ("**L**ittle-**E**ndian **B**ase **128**") is a variable-length encoding for arbitrary signed or unsigned integer quantities. The format was borrowed from the [DWARF3](#) specification. In a .dex file, LEB128 is only ever used to encode 32-bit quantities.

Each LEB128 encoded value consists of one to five bytes, which together represent a single 32-bit value. Each byte has its most significant bit set except for the final byte in the sequence, which has its most significant bit clear. The remaining seven bits of each byte are payload, with the least significant seven bits of the quantity in the first byte, the next seven in the second byte and so on. In the case of a signed LEB128 (sleb128), the most significant payload bit of the final byte in the sequence is sign-extended to produce the final value. In the unsigned case (uleb128), any bits not explicitly represented are interpreted as 0.



The variant `uleb128p1` is used to represent a signed value, where the representation is of the value *plus one* encoded as a `uleb128`. This makes the encoding of `-1` (alternatively thought of as the unsigned value `0xffffffff`) — but no other negative number — a single byte, and is useful in exactly those cases where the represented number must either be non-negative or `-1` (or `0xffffffff`), and where no other negative values are allowed (or where large unsigned values are unlikely to be needed).

Here are some examples of the formats:

| Encoded Sequence | As <code>sleb128</code> | As <code>uleb128</code> | As <code>uleb128p1</code> |
|------------------|-------------------------|-------------------------|---------------------------|
| 00 | 0 | 0 | -1 |
| 01 | 1 | 1 | 0 |
| 7f | -1 | 127 | 126 |
| 80 7f | -128 | 16256 | 16255 |

Overall File Layout

| Name | Format | Description |
|------------|-------------------|---|
| header | header_item | the header |
| string_ids | string_id_item[] | string identifiers list. These are identifiers for all the strings used by this file, either for internal naming (e.g., type descriptors) or as constant objects referred to by code. This list must be sorted by string contents, using UTF-16 code point values (not in a locale-sensitive manner). |
| type_ids | type_id_item[] | type identifiers list. These are identifiers for all types (classes, arrays, or primitive types) referred to by this file, whether defined in the file or not. This list must be sorted by <code>string_id</code> index. |
| proto_ids | proto_id_item[] | method prototype identifiers list. These are identifiers for all prototypes referred to by this file. This list must be sorted in return-type (by <code>type_id</code> index) major order, and then by arguments (also by <code>type_id</code> index). |
| field_ids | field_id_item[] | field identifiers list. These are identifiers for all fields referred to by this file, whether defined in the file or not. This list must be sorted, where the defining type (by <code>type_id</code> index) is the major order, field name (by <code>string_id</code> index) is the intermediate order, and type (by <code>type_id</code> index) is the minor order. |
| method_ids | method_id_item[] | method identifiers list. These are identifiers for all methods referred to by this file, whether defined in the file or not. This list must be sorted, where the defining type (by <code>type_id</code> index) is the major order, method name (by <code>string_id</code> index) is the intermediate |

| Name | Format | Description |
|-------------------------|--------------------------------|--|
| | | order, and method prototype (by <code>proto_id</code> index) is the minor order. |
| <code>class_defs</code> | <code>class_def_item[]</code> | class definitions list. The classes must be ordered such that a given class's superclass and implemented interfaces appear in the list earlier than the referring class. |
| <code>data</code> | <code>ubyte[]</code> | data area, containing all the support data for the tables listed above. Different items have different alignment requirements, and padding bytes are inserted before each item if necessary to achieve proper alignment. |
| <code>link_data</code> | <code>ubyte[]</code> | data used in statically linked files. The format of the data in this section is left unspecified by this document; this section is empty in unlinked files, and runtime implementations may use it as they see fit. |

Bitfield, String, and Constant Definitions

DEX_FILE_MAGIC

embedded in header_item

The constant array/string `DEX_FILE_MAGIC` is the list of bytes that must appear at the beginning of a `.dex` file in order for it to be recognized as such. The value intentionally contains a newline ("`\n`" or `0x0a`) and a null byte ("`\0`" or `0x00`) in order to help in the detection of certain forms of corruption. The value also encodes a format version number as three decimal digits, which is expected to increase monotonically over time as the format evolves.

```
ubyte[8] DEX_FILE_MAGIC = { 0x64 0x65 0x78 0x0a 0x30 0x33 0x35 0x00 }
                        = "dex\n035\0"
```

Note: At least a couple earlier versions of the format have been used in widely-available public software releases. For example, version `009` was used for the M3 releases of the Android platform (November-December 2007), and version `013` was used for the M5 releases of the Android platform (February-March 2008). In several respects, these earlier versions of the format differ significantly from the version described in this document.

ENDIAN_CONSTANT and REVERSE_ENDIAN_CONSTANT

embedded in header_item

The constant `ENDIAN_CONSTANT` is used to indicate the endianness of the file in which it is found. Although the standard `.dex` format is little-endian, implementations may choose to perform byte-swapping. Should an implementation come across a header whose `endian_tag` is `REVERSE_ENDIAN_CONSTANT` instead of `ENDIAN_CONSTANT`, it would know that the file has been byte-swapped from the expected form.

```
uint ENDIAN_CONSTANT = 0x12345678;
uint REVERSE_ENDIAN_CONSTANT = 0x78563412;
```

NO_INDEX

embedded in class_def_item and debug_info_item

The constant NO_INDEX is used to indicate that an index value is absent.

Note: This value isn't defined to be 0, because that is in fact typically a valid index.

Also Note: The chosen value for NO_INDEX is representable as a single byte in the uleb128p1 encoding.

```
uint NO_INDEX = 0xffffffff;    // == -1 if treated as a signed int
```

access_flags Definitions

embedded in class_def_item, field_item, method_item, and InnerClass

Bitfields of these flags are used to indicate the accessibility and overall properties of classes and class members.

| Name | Value | For Classes (and InnerClass annotations) | For Fields | For Methods |
|---------------|-------|---|---|--|
| ACC_PUBLIC | 0x1 | public: visible everywhere | public: visible everywhere | public: visible everywhere |
| ACC_PRIVATE | 0x2 | * private: only visible to defining class | private: only visible to defining class | private: only visible to defining class |
| ACC_PROTECTED | 0x4 | * protected: visible to package and subclasses | protected: visible to package and subclasses | protected: visible to package and subclasses |
| ACC_STATIC | 0x8 | * static: is not constructed with an outer this reference | static: global to defining class | static: does not take a this argument |
| ACC_FINAL | 0x10 | final: not subclassable | final: immutable after construction | final: not overridable |
| (unused) | 0x20 | | | |
| ACC_VOLATILE | 0x40 | | volatile: special access rules to help with thread safety | |
| ACC_BRIDGE | 0x40 | | | bridge method, added automatically by compiler as a type-safe bridge |
| ACC_TRANSIENT | 0x80 | | transient: not to be saved by default serialization | |

| Name | Value | For Classes (and InnerClass annotations) | For Fields | For Methods |
|-----------------|---------|--|--|--|
| ACC_VARARGS | 0x80 | | | last argument should be treated as a "rest" argument by compiler |
| ACC_NATIVE | 0x100 | | | native: implemented in native code |
| ACC_INTERFACE | 0x200 | interface: multiply-implementable abstract class | | |
| ACC_ABSTRACT | 0x400 | abstract: not directly instantiable | | abstract: unimplemented by this class |
| ACC_STRICT | 0x800 | | | strictfp: strict rules for floating-point arithmetic |
| ACC_SYNTHETIC | 0x1000 | not directly defined in source code | not directly defined in source code | not directly defined in source code |
| ACC_ANNOTATION | 0x2000 | declared as an annotation class | | |
| ACC_ENUM | 0x4000 | declared as an enumerated type | declared as an enumerated value | |
| (unused) | 0x8000 | | | |
| ACC_CONSTRUCTOR | 0x10000 | | | constructor method (class or instance initializer) |

* Only allowed on for `InnerClass` annotations, and must not ever be on in a `class_def_item`.

MUTF-8 (Modified UTF-8) Encoding

As a concession to easier legacy support, the `.dex` format encodes its string data in a de facto standard modified UTF-8 form, hereafter referred to as MUTF-8. This form is identical to standard UTF-8, except:

- Only the one-, two-, and three-byte encodings are used.
- Code points in the range `U+10000 ... U+10ffff` are encoded as a surrogate pair, each of which is represented as a three-byte encoded value.
- The code point `U+0000` is encoded in two-byte form.
- A plain null byte (value 0) indicates the end of a string, as is the standard C language interpretation.

The first two items above can be summarized as: MUTF-8 is an encoding format for UTF-16, instead of being a more direct encoding format for Unicode characters.

The final two items above make it simultaneously possible to include the code point `U+0000` in a string *and* still manipulate it as a C-style null-terminated string.

However, the special encoding of `U+0000` means that, unlike normal UTF-8, the result of calling the standard C function `strcmp()` on a pair of MUTF-8 strings does not always indicate the properly signed result of comparison of *unequal* strings. When ordering (not just equality) is a concern, the most straightforward way to compare MUTF-8 strings is to decode them character by character, and compare the decoded values. (However, more clever implementations are also possible.)

Please refer to [The Unicode Standard](#) for further information about character encoding. UTF-8 is actually closer to the (relatively less well-known) encoding [CESU-8](#) than to UTF-8 per se.

encoded_value Encoding

embedded in annotation_element and encoded_array_item

An encoded_value is an encoded piece of (nearly) arbitrary hierarchically structured data. The encoding is meant to be both compact and straightforward to parse.

| Name | Format | Description |
|--|----------|--|
| <code>(value_arg << 5) value_type</code> | ubyte | byte indicating the type of the immediately subsequent value along with an optional clarifying argument in the high-order three bits. See below for the various value definitions. In most cases, value_arg encodes the length of the immediately-subsequent value in bytes, as (size - 1), e.g., 0 means that the value requires one byte, and 7 means it requires eight bytes; however, there are exceptions as noted below. |
| value | ubyte[] | bytes representing the value, variable in length and interpreted differently for different value_type bytes, though always little-endian. See the various value definitions below for details. |

Value Formats

| Type Name | value_type | value_arg Format | value Format | Description |
|--------------|------------|-------------------|--------------|--|
| VALUE_BYTE | 0x00 | (none; must be 0) | ubyte[1] | signed one-byte integer value |
| VALUE_SHORT | 0x02 | size - 1 (0..1) | ubyte[size] | signed two-byte integer value, sign-extended |
| VALUE_CHAR | 0x03 | size - 1 (0..1) | ubyte[size] | unsigned two-byte integer value, zero-extended |
| VALUE_INT | 0x04 | size - 1 (0..3) | ubyte[size] | signed four-byte integer value, sign-extended |
| VALUE_LONG | 0x06 | size - 1 (0..7) | ubyte[size] | signed eight-byte integer value, sign-extended |
| VALUE_FLOAT | 0x10 | size - 1 (0..3) | ubyte[size] | four-byte bit pattern, zero-extended to the right, and interpreted as an IEEE754 32-bit floating point value |
| VALUE_DOUBLE | 0x11 | size - 1 (0..7) | ubyte[size] | eight-byte bit pattern, zero-extended to the right, and interpreted as an IEEE754 64-bit floating point value |
| VALUE_STRING | 0x17 | size - 1 (0..3) | ubyte[size] | unsigned (zero-extended) four-byte integer value, interpreted as an index into the string_ids section and representing a string value |
| VALUE_TYPE | 0x18 | size - 1 (0..3) | ubyte[size] | unsigned (zero-extended) four-byte integer value, interpreted as an index into the type_ids section and representing a reflective type/class value |

| Type Name | value_type | value_arg Format | value Format | Description |
|------------------|------------|-------------------|--------------------|--|
| VALUE_FIELD | 0x19 | size - 1 (0..3) | ubyte[size] | unsigned (zero-extended) four-byte integer value, interpreted as an index into the field_ids section and representing a reflective field value |
| VALUE_METHOD | 0x1a | size - 1 (0..3) | ubyte[size] | unsigned (zero-extended) four-byte integer value, interpreted as an index into the method_ids section and representing a reflective method value |
| VALUE_ENUM | 0x1b | size - 1 (0..3) | ubyte[size] | unsigned (zero-extended) four-byte integer value, interpreted as an index into the field_ids section and representing the value of an enumerated type constant |
| VALUE_ARRAY | 0x1c | (none; must be 0) | encoded_array | an array of values, in the format specified by "encoded_array Format" below. The size of the value is implicit in the encoding. |
| VALUE_ANNOTATION | 0x1d | (none; must be 0) | encoded_annotation | a sub-annotation, in the format specified by "encoded_annotation Format" below. The size of the value is implicit in the encoding. |
| VALUE_NULL | 0x1e | (none; must be 0) | (none) | null reference value |
| VALUE_BOOLEAN | 0x1f | boolean (0..1) | (none) | one-bit value; 0 for false and 1 for true. The bit is represented in the value_arg. |

encoded_array Format

| Name | Format | Description |
|--------|---------------------|---|
| size | uleb128 | number of elements in the array |
| values | encoded_value[size] | a series of size encoded_value byte sequences in the format specified by this section, concatenated sequentially. |

encoded_annotation Format

| Name | Format | Description |
|----------|--------------------------|---|
| type_idx | uleb128 | type of the annotation. This must be a class (not array or primitive) type. |
| size | uleb128 | number of name-value mappings in this annotation |
| elements | annotation_element[size] | elements of the annotataion, represented directly in-line (not as offsets). Elements must be sorted in increasing order by string_id index. |

annotation_element Format

| Name | Format | Description |
|----------|---------------|--|
| name_idx | uleb128 | element name, represented as an index into the <code>string_ids</code> section. The string must conform to the syntax for <i>MemberName</i> , defined above. |
| value | encoded_value | element value |

String Syntax

There are several kinds of item in a `.dex` file which ultimately refer to a string. The following BNF-style definitions indicate the acceptable syntax for these strings.

SimpleName

A *SimpleName* is the basis for the syntax of the names of other things. The `.dex` format allows a fair amount of latitude here (much more than most common source languages). In brief, a simple name may consist of any low-ASCII alphabetic character or digit, a few specific low-ASCII symbols, and most non-ASCII code points that are not control, space, or special characters. Note that surrogate code points (in the range `U+d800 ... U+ffff`) are not considered valid name characters, per se, but Unicode supplemental characters *are* valid (which are represented by the final alternative of the rule for *SimpleNameChar*), and they should be represented in a file as pairs of surrogate code points in the UTF-8 encoding.

```

SimpleName →
    SimpleNameChar (SimpleNameChar)*
SimpleNameChar →
    'A' ... 'Z'
  | 'a' ... 'z'
  | '0' ... '9'
  | '$'
  | '-'
  | '_'
  | U+00a1 ... U+1fff
  | U+2010 ... U+2027
  | U+2030 ... U+d7ff
  | U+e000 ... U+ffef
  | U+10000 ... U+10ffff

```

MemberName

used by *field_id_item* and *method_id_item*

A *MemberName* is the name of a member of a class, members being fields, methods, and inner classes.

```

MemberName →
    SimpleName
  | '<' SimpleName '>'

```


FullClassName

A *FullClassName* is a fully-qualified class name, including an optional package specifier followed by a required name.

```
FullClassName →  
    OptionalPackagePrefix SimpleName  
  
OptionalPackagePrefix →  
    (SimpleName ' / ')*
```

TypeDescriptor

used by type_id_item

A *TypeDescriptor* is the representation of any type, including primitives, classes, arrays, and void. See below for the meaning of the various versions.

```
TypeDescriptor →  
    'V'  
    | FieldTypeDescriptor  
  
FieldTypeDescriptor →  
    NonArrayFieldTypeDescriptor  
    | ( '[' * 1...255 ) NonArrayFieldTypeDescriptor  
  
NonArrayFieldTypeDescriptor →  
    'Z'  
    | 'B'  
    | 'S'  
    | 'C'  
    | 'I'  
    | 'J'  
    | 'F'  
    | 'D'  
    | 'L' FullClassName ' ; '
```

ShortyDescriptor

used by proto_id_item

A *ShortyDescriptor* is the short form representation of a method prototype, including return and parameter types, except that there is no distinction between various reference (class or array) types. Instead, all reference types are represented by a single 'L' character.

```
ShortyDescriptor →  
    ShortyReturnType (ShortyFieldType)*  
  
ShortyReturnType →  
    'V'  
    | ShortyFieldType  
  
ShortyFieldType →  
    'Z'  
    | 'B'  
    | 'S'  
    | 'C'
```

```
| 'I'  
| 'J'  
| 'F'  
| 'D'  
| 'L'
```

TypeDescriptor Semantics

This is the meaning of each of the variants of *TypeDescriptor*.

| Syntax | Meaning |
|---------------------------------|--|
| V | void; only valid for return types |
| Z | boolean |
| B | byte |
| S | short |
| C | char |
| I | int |
| J | long |
| F | float |
| D | double |
| L <i>fully/qualified/Name</i> ; | the class <i>fully.qualified.Name</i> |
| [<i>descriptor</i> | array of <i>descriptor</i> , usable recursively for arrays-of-arrays, though it is invalid to have more than 255 dimensions. |

Items and Related Structures

This section includes definitions for each of the top-level items that may appear in a `.dex` file.

header_item

appears in the header section

alignment: 4 bytes

| Name | Format | Description |
|----------|---------------------------|--|
| magic | ubyte[8] = DEX_FILE_MAGIC | magic value. See discussion above under "DEX_FILE_MAGIC" for more details. |
| checksum | uint | adler32 checksum of the rest of the file (everything but magic and this field); used to detect file corruption |

map_list

appears in the data section

referenced from header_item

alignment: 4 bytes

This is a list of the entire contents of a file, in order. It contains some redundancy with respect to the header_item but is intended to be an easy form to use to iterate over an entire file. A given type may appear at most once in a map, but there is no restriction on what order types may appear in, other than the restrictions implied by the rest of the format (e.g., a header section must appear first, followed by a string_ids section, etc.). Additionally, the map entries must be ordered by initial offset and must not overlap.

| Name | Format | Description |
|------|----------------|------------------------------|
| size | uint | size of the list, in entries |
| list | map_item[size] | elements of the list |

map_item Format

| Name | Format | Description |
|--------|--------|--|
| type | ushort | type of the items; see table below |
| unused | ushort | (unused) |
| size | uint | count of the number of items to be found at the indicated offset |
| offset | uint | offset from the start of the file to the items in question |

Type Codes

| Item Type | Constant | Value | Item Size In Bytes |
|----------------|---------------------|--------|----------------------------|
| header_item | TYPE_HEADER_ITEM | 0x0000 | 0x70 |
| string_id_item | TYPE_STRING_ID_ITEM | 0x0001 | 0x04 |
| type_id_item | TYPE_TYPE_ID_ITEM | 0x0002 | 0x04 |
| proto_id_item | TYPE_PROTO_ID_ITEM | 0x0003 | 0x0c |
| field_id_item | TYPE_FIELD_ID_ITEM | 0x0004 | 0x08 |
| method_id_item | TYPE_METHOD_ID_ITEM | 0x0005 | 0x08 |
| class_def_item | TYPE_CLASS_DEF_ITEM | 0x0006 | 0x20 |
| map_list | TYPE_MAP_LIST | 0x1000 | 4 + (item.size * 12) |
| type_list | TYPE_TYPE_LIST | 0x1001 | 4 + (item.size * 2) |

| Item Type | Constant | Value | Item Size In Bytes |
|----------------------------|---------------------------------|--------|-----------------------------|
| annotation_set_ref_list | TYPE_ANNOTATION_SET_REF_LIST | 0x1002 | 4 + (item.size * 4) |
| annotation_set_item | TYPE_ANNOTATION_SET_ITEM | 0x1003 | 4 + (item.size * 4) |
| class_data_item | TYPE_CLASS_DATA_ITEM | 0x2000 | <i>implicit; must parse</i> |
| code_item | TYPE_CODE_ITEM | 0x2001 | <i>implicit; must parse</i> |
| string_data_item | TYPE_STRING_DATA_ITEM | 0x2002 | <i>implicit; must parse</i> |
| debug_info_item | TYPE_DEBUG_INFO_ITEM | 0x2003 | <i>implicit; must parse</i> |
| annotation_item | TYPE_ANNOTATION_ITEM | 0x2004 | <i>implicit; must parse</i> |
| encoded_array_item | TYPE_ENCODED_ARRAY_ITEM | 0x2005 | <i>implicit; must parse</i> |
| annotations_directory_item | TYPE_ANNOTATIONS_DIRECTORY_ITEM | 0x2006 | <i>implicit; must parse</i> |

string_id_item

appears in the string_ids section

alignment: 4 bytes

| Name | Format | Description |
|-----------------|--------|---|
| string_data_off | uint | offset from the start of the file to the string data for this item. The offset should be to a location in the data section, and the data should be in the format specified by "string_data_item" below. There is no alignment requirement for the offset. |

string_data_item

appears in the data section

alignment: none (byte-aligned)

| Name | Format | Description |
|------------|----------|--|
| utf16_size | uleb128 | size of this string, in UTF-16 code units (which is the "string length" in many systems). That is, this is the decoded length of the string. (The encoded length is implied by the position of the 0 byte.) |
| data | ubyte[] | <p>a series of MUTF-8 code units (a.k.a. octets, a.k.a. bytes) followed by a byte of value 0. See "MUTF-8 (Modified UTF-8) Encoding" above for details and discussion about the data format.</p> <p>Note: It is acceptable to have a string which includes (the encoded form of) UTF-16 surrogate code units (that is, U+d800 ... U+ffff) either in isolation or out-of-order with respect to the usual encoding of Unicode into UTF-16. It is up to higher-level uses of strings to reject such invalid encodings, if appropriate.</p> |

type_id_item

appears in the `type_ids` section

alignment: 4 bytes

| Name | Format | Description |
|----------------|--------|--|
| descriptor_idx | uint | index into the <code>string_ids</code> list for the descriptor string of this type. The string must conform to the syntax for <i>TypeDescriptor</i> , defined above. |

proto_id_item

appears in the `proto_ids` section

alignment: 4 bytes

| Name | Format | Description |
|-----------------|--------|--|
| shorty_idx | uint | index into the <code>string_ids</code> list for the short-form descriptor string of this prototype. The string must conform to the syntax for <i>ShortyDescriptor</i> , defined above, and must correspond to the return type and parameters of this item. |
| return_type_idx | uint | index into the <code>type_ids</code> list for the return type of this prototype |
| parameters_off | uint | offset from the start of the file to the list of parameter types for this prototype, or 0 if this prototype has no parameters. This offset, if non-zero, should be in the data section, and the data there should be in the format specified by "type_list" below. Additionally, there should be no reference to the type <code>void</code> in the list. |

field_id_item

appears in the `field_ids` section

alignment: 4 bytes

| Name | Format | Description |
|-----------|--------|--|
| class_idx | ushort | index into the <code>type_ids</code> list for the definer of this field. This must be a class type, and not an array or primitive type. |
| type_idx | ushort | index into the <code>type_ids</code> list for the type of this field |
| name_idx | uint | index into the <code>string_ids</code> list for the name of this field. The string must conform to the syntax for <i>MemberName</i> , defined above. |

method_id_item

appears in the `method_ids` section

alignment: 4 bytes

| Name | Format | Description |
|------------------------|---------------------|---|
| <code>class_idx</code> | <code>ushort</code> | index into the <code>type_ids</code> list for the definer of this method. This must be a class or array type, and not a primitive type. |
| <code>proto_idx</code> | <code>ushort</code> | index into the <code>proto_ids</code> list for the prototype of this method |
| <code>name_idx</code> | <code>uint</code> | index into the <code>string_ids</code> list for the name of this method. The string must conform to the syntax for <i>MemberName</i> , defined above. |

class_def_item

appears in the `class_defs` section

alignment: 4 bytes

| Name | Format | Description |
|--------------------------------|-------------------|--|
| <code>class_idx</code> | <code>uint</code> | index into the <code>type_ids</code> list for this class. This must be a class type, and not an array or primitive type. |
| <code>access_flags</code> | <code>uint</code> | access flags for the class (public, final, etc.). See "access_flags Definitions" for details. |
| <code>superclass_idx</code> | <code>uint</code> | index into the <code>type_ids</code> list for the superclass, or the constant value <code>NO_INDEX</code> if this class has no superclass (i.e., it is a root class such as <code>Object</code>). If present, this must be a class type, and not an array or primitive type. |
| <code>interfaces_off</code> | <code>uint</code> | offset from the start of the file to the list of interfaces, or 0 if there are none. This offset should be in the data section, and the data there should be in the format specified by "type_list" below. Each of the elements of the list must be a class type (not an array or primitive type), and there must not be any duplicates. |
| <code>source_file_idx</code> | <code>uint</code> | index into the <code>string_ids</code> list for the name of the file containing the original source for (at least most of) this class, or the special value <code>NO_INDEX</code> to represent a lack of this information. The <code>debug_info_item</code> of any given method may override this source file, but the expectation is that most classes will only come from one source file. |
| <code>annotations_off</code> | <code>uint</code> | offset from the start of the file to the annotations structure for this class, or 0 if there are no annotations on this class. This offset, if non-zero, should be in the data section, and the data there should be in the format specified by "annotations_directory_item" below, with all items referring to this class as the definer. |
| <code>class_data_off</code> | <code>uint</code> | offset from the start of the file to the associated class data for this item, or 0 if there is no class data for this class. (This may be the case, for example, if this class is a marker interface.) The offset, if non-zero, should be in the data section, and the data there should be in the format specified by "class_data_item" below, with all items referring to this class as the definer. |
| <code>static_values_off</code> | <code>uint</code> | offset from the start of the file to the list of initial values for static fields, or 0 if there are none (and all <code>static</code> fields are to be initialized with 0 or null). This offset should be in the data section, and the data there should be in the format specified by |

| Name | Format | Description |
|------|--------|--|
| | | "encoded_array_item" below. The size of the array must be no larger than the number of <code>static</code> fields declared by this class, and the elements correspond to the <code>static</code> fields in the same order as declared in the corresponding <code>field_list</code> . The type of each array element must match the declared type of its corresponding field. If there are fewer elements in the array than there are <code>static</code> fields, then the leftover fields are initialized with a type-appropriate 0 or null. |

class_data_item

referenced from `class_def_item`

appears in the `data` section

alignment: none (byte-aligned)

| Name | Format | Description |
|-----------------------------------|---|--|
| <code>static_fields_size</code> | <code>uleb128</code> | the number of static fields defined in this item |
| <code>instance_fields_size</code> | <code>uleb128</code> | the number of instance fields defined in this item |
| <code>direct_methods_size</code> | <code>uleb128</code> | the number of direct methods defined in this item |
| <code>virtual_methods_size</code> | <code>uleb128</code> | the number of virtual methods defined in this item |
| <code>static_fields</code> | <code>encoded_field[static_fields_size]</code> | the defined static fields, represented as a sequence of encoded elements. The fields must be sorted by <code>field_idx</code> in increasing order. |
| <code>instance_fields</code> | <code>encoded_field[instance_fields_size]</code> | the defined instance fields, represented as a sequence of encoded elements. The fields must be sorted by <code>field_idx</code> in increasing order. |
| <code>direct_methods</code> | <code>encoded_method[direct_methods_size]</code> | the defined direct (any of <code>static</code> , <code>private</code> , or <code>constructor</code>) methods, represented as a sequence of encoded elements. The methods must be sorted by <code>method_idx</code> in increasing order. |
| <code>virtual_methods</code> | <code>encoded_method[virtual_methods_size]</code> | the defined virtual (none of <code>static</code> , <code>private</code> , or <code>constructor</code>) methods, represented as a sequence of encoded elements. This list should <i>not</i> include inherited methods unless overridden by the class that this item represents. The methods must be sorted by <code>method_idx</code> in increasing order. |

Note: All elements' `field_ids` and `method_ids` must refer to the same defining class.

encoded_field Format

| Name | Format | Description |
|-----------------------------|----------------------|---|
| <code>field_idx_diff</code> | <code>uleb128</code> | index into the <code>field_ids</code> list for the identity of this field (includes the name and descriptor), represented as a difference from the index of previous element in the list. The index of the first element in a list is represented directly. |
| <code>access_flags</code> | <code>uleb128</code> | access flags for the field (<code>public</code> , <code>final</code> , etc.). See "access_flags Definitions" for details. |

encoded_method Format

| Name | Format | Description |
|------------------------------|----------------------|---|
| <code>method_idx_diff</code> | <code>uleb128</code> | index into the <code>method_ids</code> list for the identity of this method (includes the name and descriptor), represented as a difference from the index of previous element in the list. The index of the first element in a list is represented directly. |
| <code>access_flags</code> | <code>uleb128</code> | access flags for the method (<code>public</code> , <code>final</code> , etc.). See "access_flags Definitions" for details. |
| <code>code_off</code> | <code>uleb128</code> | offset from the start of the file to the code structure for this method, or 0 if this method is either abstract or native. The offset should be to a location in the data section. The format of the data is specified by "code_item" below. |

type_list

referenced from `class_def_item` and `proto_id_item`

appears in the data section

alignment: 4 bytes

| Name | Format | Description |
|-------------------|------------------------------|------------------------------|
| <code>size</code> | <code>uint</code> | size of the list, in entries |
| <code>list</code> | <code>type_item[size]</code> | elements of the list |

type_item Format

| Name | Format | Description |
|-----------------------|---------------------|---|
| <code>type_idx</code> | <code>ushort</code> | index into the <code>type_ids</code> list |

code_item

referenced from *method_item*

appears in the *data* section

alignment: 4 bytes

| Name | Format | Description |
|----------------|---------------------------------------|---|
| registers_size | ushort | the number of registers used by this code |
| ins_size | ushort | the number of words of incoming arguments to the method that this code is for |
| outs_size | ushort | the number of words of outgoing argument space required by this code for method invocation |
| tries_size | ushort | the number of <i>try_items</i> for this instance. If non-zero, then these appear as the <i>tries</i> array just after the <i>insns</i> in this instance. |
| debug_info_off | uint | offset from the start of the file to the debug info (line numbers + local variable info) sequence for this code, or 0 if there simply is no information. The offset, if non-zero, should be to a location in the data section. The format of the data is specified by "debug_info_item" below. |
| insns_size | uint | size of the instructions list, in 16-bit code units |
| insns | ushort[insns_size] | actual array of bytecode. The format of code in an <i>insns</i> array is specified by the companion document " Bytecode for the Dalvik VM ". Note that though this is defined as an array of <i>ushort</i> , there are some internal structures that prefer four-byte alignment. Also, if this happens to be in an endian-swapped file, then the swapping is <i>only</i> done on individual <i>ushorts</i> and not on the larger internal structures. |
| padding | ushort (optional) = 0 | two bytes of padding to make <i>tries</i> four-byte aligned. This element is only present if <i>tries_size</i> is non-zero and <i>insns_size</i> is odd. |
| tries | try_item[tries_size] (optional) | array indicating where in the code exceptions may be caught and how to handle them. Elements of the array must be non-overlapping in range and in order from low to high address. This element is only present if <i>tries_size</i> is non-zero. |
| handlers | encoded_catch_handler_list (optional) | bytes representing a list of lists of catch types and associated handler addresses. Each <i>try_item</i> has a byte-wise offset into this structure. This element is only present if <i>tries_size</i> is non-zero. |

try_item Format

| Name | Format | Description |
|-------------|--------|--|
| start_addr | uint | start address of the block of code covered by this entry. The address is a count of 16-bit code units to the start of the first covered instruction. |
| insn_count | ushort | number of 16-bit code units covered by this entry. The last code unit covered (inclusive) is <i>start_addr</i> + <i>insn_count</i> - 1. |
| handler_off | ushort | offset in bytes from the start of the associated encoded handler data to the <i>catch_handler_item</i> for this entry |

encoded_catch_handler_list Format

| Name | Format | Description |
|------|--------------------------------------|--|
| size | uleb128 | size of this list, in entries |
| list | encoded_catch_handler[handlers_size] | actual list of handler lists, represented directly (not as offsets), and concatenated sequentially |

encoded_catch_handler Format

| Name | Format | Description |
|----------------|-----------------------------------|--|
| size | sleb128 | number of catch types in this list. If non-positive, then this is the negative of the number of catch types, and the catches are followed by a catch-all handler. For example: A size of 0 means that there is a catch-all but no explicitly typed catches. A size of 2 means that there are two explicitly typed catches and no catch-all. And a size of -1 means that there is one typed catch along with a catch-all. |
| handlers | encoded_type_addr_pair[abs(size)] | stream of abs(size) encoded items, one for each caught type, in the order that the types should be tested. |
| catch_all_addr | uleb128 (optional) | bytecode address of the catch-all handler. This element is only present if size is non-positive. |

encoded_type_addr_pair Format

| Name | Format | Description |
|----------|---------|---|
| type_idx | uleb128 | index into the type_ids list for the type of the exception to catch |
| addr | uleb128 | bytecode address of the associated exception handler |

debug_info_item

referenced from code_item

appears in the data section

alignment: none (byte-aligned)

Each debug_info_item defines a DWARF3-inspired byte-coded state machine that, when interpreted, emits the positions table and (potentially) the local variable information for a code_item. The sequence begins with a variable-length header (the length of which depends on the number of method parameters), is followed by the state machine bytecodes, and ends with an DBG_END_SEQUENCE byte.

The state machine consists of five registers. The address register represents the instruction offset in the

associated `insns_item` in 16-bit code units. The address register starts at 0 at the beginning of each `debug_info` sequence and may only monotonically increase. The `line` register represents what source line number should be associated with the next positions table entry emitted by the state machine. It is initialized in the sequence header, and may change in positive or negative directions but must never be less than 1. The `source_file` register represents the source file that the line number entries refer to. It is initialized to the value of `source_file_idx` in `class_def_item`. The other two variables, `prologue_end` and `epilogue_begin`, are boolean flags (initialized to `false`) that indicate whether the next position emitted should be considered a method prologue or epilogue. The state machine must also track the name and type of the last local variable live in each register for the `DBG_RESTART_LOCAL` code.

The header is as follows:

| Name | Format | Description |
|------------------------------|---|--|
| <code>line_start</code> | <code>uleb128</code> | the initial value for the state machine's line register. Does not represent an actual positions entry. |
| <code>parameters_size</code> | <code>uleb128</code> | the number of parameter names that are encoded. There should be one per method parameter, excluding an instance method's <code>this</code> , if any. |
| <code>parameter_names</code> | <code>uleb128p1[parameters_size]</code> | string index of the method parameter name. An encoded value of <code>NO_INDEX</code> indicates that no name is available for the associated parameter. The type descriptor and signature are implied from the method descriptor and signature. |

The byte code values are as follows:

| Name | Value | Format | Arguments | Description |
|---------------------------------------|-------------------|---|---|---|
| <code>DBG_END_SEQUENCE</code> | <code>0x00</code> | | <i>(none)</i> | terminates a debug info sequence <code>code_item</code> |
| <code>DBG_ADVANCE_PC</code> | <code>0x01</code> | <code>uleb128</code> <code>addr_diff</code> | <code>addr_diff</code> : amount to add to address register | advances the address register emitting a positions entry |
| <code>DBG_ADVANCE_LINE</code> | <code>0x02</code> | <code>sleb128</code> <code>line_diff</code> | <code>line_diff</code> : amount to change line register by | advances the line register with a positions entry |
| <code>DBG_START_LOCAL</code> | <code>0x03</code> | <code>uleb128</code> <code>register_num</code> <code>uleb128p1</code> <code>name_idx</code> <code>uleb128p1</code> <code>type_idx</code> | <code>register_num</code> : register that will contain local <code>name_idx</code> : string index of the name <code>type_idx</code> : type index of the type | introduces a local variable at the current address. Either <code>name_idx</code> or <code>type_idx</code> may be <code>NO_INDEX</code> to indicate that that value is unknown. |
| <code>DBG_START_LOCAL_EXTENDED</code> | <code>0x04</code> | <code>uleb128</code> <code>register_num</code> <code>uleb128p1</code> <code>name_idx</code> <code>uleb128p1</code> <code>type_idx</code> <code>uleb128p1</code> <code>sig_idx</code> | <code>register_num</code> : register that will contain local <code>name_idx</code> : string index of the name <code>type_idx</code> : type index of the type <code>sig_idx</code> : string index of the type signature | introduces a local with a type at the current address. Any of <code>name_idx</code> , <code>type_idx</code> , or <code>sig_idx</code> may be <code>NO_INDEX</code> to indicate that that value is unknown. (If <code>sig_idx</code> is <code>-1</code> , same data could be represented efficiently using the opcode <code>DBG_START_LOCAL</code> .) Note: See the discussion under <code>"dalvik.annotation.S"</code> below for caveats about handling signatures. |

| Name | Value | Format | Arguments | Description |
|------------------------|-------------|----------------------|---|--|
| DBG_END_LOCAL | 0x05 | uleb128 register_num | register_num: register that contained local | marks a currently-live local variable out of scope at the current address. |
| DBG_RESTART_LOCAL | 0x06 | uleb128 register_num | register_num: register to restart | re-introduces a local variable at the current address. The name and register are the same as the last local that used the specified register. |
| DBG_SET_PROLOGUE_END | 0x07 | | (none) | sets the prologue_end stack machine register, indicating that the next position entry that is added should be considered the end of a method prologue (an appropriate place for a method breakpoint). The prologue_end register is cleared by any special (>= 0x0a) opcode. |
| DBG_SET_EPILOGUE_BEGIN | 0x08 | | (none) | sets the epilogue_begin stack machine register, indicating the position entry that is added should be considered the beginning of an epilogue (an appropriate place for execution before method exit). The epilogue_begin register is cleared by any special (>= 0x0a) opcode. |
| DBG_SET_FILE | 0x09 | uleb128pl | name_idx: string index of source file name; NO_INDEX if unknown | indicates that all subsequent line entries make reference to this source file name, instead of the default name specified in code_item. |
| Special Opcodes | 0x0a...0xff | | (none) | advances the line and address registers, emits a new position table entry, clears prologue_end and epilogue_begin. See below for description. |

Special Opcodes

Opcodes with values between 0x0a and 0xff (inclusive) move both the line and address registers by a small amount and then emit a new position table entry. The formula for the increments are as follows:

```

DBG_FIRST_SPECIAL = 0x0a // the smallest special opcode
DBG_LINE_BASE     = -4    // the smallest line number increment
DBG_LINE_RANGE    = 15    // the number of line increments represented

adjusted_opcode = opcode - DBG_FIRST_SPECIAL

line += DBG_LINE_BASE + (adjusted_opcode % DBG_LINE_RANGE)
address += (adjusted_opcode / DBG_LINE_RANGE)

```

annotations_directory_item

referenced from *class_def_item*

appears in the *data* section

alignment: 4 bytes

| Name | Format | Description |
|--------------------------|---|---|
| class_annotations_off | uint | offset from the start of the file to the annotations made directly on the class, or 0 if the class has no direct annotations. The offset, if non-zero, should be to a location in the data section. The format of the data is specified by "annotation_set_item" below. |
| fields_size | uint | count of fields annotated by this item |
| annotated_methods_off | uint | count of methods annotated by this item |
| annotated_parameters_off | uint | count of method parameter lists annotated by this item |
| field_annotations | field_annotation[fields_size] <i>(optional)</i> | list of associated field annotations. The elements of the list must be sorted in increasing order, by field_idx. |
| method_annotations | method_annotation[methods_size] <i>(optional)</i> | list of associated method annotations. The elements of the list must be sorted in increasing order, by method_idx. |
| parameter_annotations | parameter_annotation[parameters_size] <i>(optional)</i> | list of associated method parameter annotations. The elements of the list must be sorted in increasing order, by method_idx. |

Note: All elements' field_ids and method_ids must refer to the same defining class.

field_annotation Format

| Name | Format | Description |
|-----------------|--------|---|
| field_idx | uint | index into the field_ids list for the identity of the field being annotated |
| annotations_off | uint | offset from the start of the file to the list of annotations for the field. The offset should be to a location in the data section. The format of the data is specified by "annotation_set_item" below. |

method_annotation Format

| Name | Format | Description |
|------------|--------|---|
| method_idx | uint | index into the method_ids list for the identity of the method being annotated |

| Name | Format | Description |
|-----------------|--------|--|
| annotations_off | uint | offset from the start of the file to the list of annotations for the method. The offset should be to a location in the data section. The format of the data is specified by "annotation_set_item" below. |

parameter_annotation Format

| Name | Format | Description |
|-----------------|--------|---|
| method_idx | uint | index into the method_ids list for the identity of the method whose parameters are being annotated |
| annotations_off | uint | offset from the start of the file to the list of annotations for the method parameters. The offset should be to a location in the data section. The format of the data is specified by "annotation_set_ref_list" below. |

annotation_set_ref_list

referenced from parameter_annotations_item

appears in the data section

alignment: 4 bytes

| Name | Format | Description |
|------|-------------------------------|------------------------------|
| size | uint | size of the list, in entries |
| list | annotation_set_ref_item[size] | elements of the list |

annotation_set_ref_item Format

| Name | Format | Description |
|-----------------|--------|---|
| annotations_off | uint | offset from the start of the file to the referenced annotation set or 0 if there are no annotations for this element. The offset, if non-zero, should be to a location in the data section. The format of the data is specified by "annotation_set_item" below. |

annotation_set_item

referenced from annotations_directory_item, field_annotations_item,

method_annotations_item, and annotation_set_ref_item

appears in the data section

alignment: 4 bytes

| Name | Format | Description |
|---------|---------------------------|--|
| size | uint | size of the set, in entries |
| entries | annotation_off_item[size] | elements of the set. The elements must be sorted in increasing order, by type_idx. |

annotation_off_item Format

| Name | Format | Description |
|----------------|--------|--|
| annotation_off | uint | offset from the start of the file to an annotation. The offset should be to a location in the data section, and the format of the data at that location is specified by "annotation_item" below. |

annotation_item

referenced from annotation_set_item

appears in the data section

alignment: none (byte-aligned)

| Name | Format | Description |
|------------|--------------------|---|
| visibility | ubyte | intended visibility of this annotation (see below) |
| annotation | encoded_annotation | encoded annotation contents, in the format described by "encoded_annotation Format" under "encoded_value Encoding" above. |

Visibility values

These are the options for the visibility field in an annotation_item:

| Name | Value | Description |
|--------------------|-------|--|
| VISIBILITY_BUILD | 0x00 | intended only to be visible at build time (e.g., during compilation of other code) |
| VISIBILITY_RUNTIME | 0x01 | intended to visible at runtime |
| VISIBILITY_SYSTEM | 0x02 | intended to visible at runtime, but only to the underlying system (and not to regular user code) |

encoded_array_item

referenced from *class_def_item*

appears in the *data* section

alignment: none (byte-aligned)

| Name | Format | Description |
|-------|---------------|---|
| value | encoded_array | bytes representing the encoded array value, in the format specified by "encoded_array Format" under "encoded_value Encoding" above. |

System Annotations

System annotations are used to represent various pieces of reflective information about classes (and methods and fields). This information is generally only accessed indirectly by client (non-system) code.

System annotations are represented in .dex files as annotations with visibility set to `VISIBILITY_SYSTEM`.

dalvik.annotation.AnnotationDefault

appears on methods in annotation interfaces

An `AnnotationDefault` annotation is attached to each annotation interface which wishes to indicate default bindings.

| Name | Format | Description |
|-------|------------|--|
| value | Annotation | the default bindings for this annotation, represented as an annotation of this type. The annotation need not include all names defined by the annotation; missing names simply do not have defaults. |

dalvik.annotation.EnclosingClass

appears on classes

An `EnclosingClass` annotation is attached to each class which is either defined as a member of another class, per se, or is anonymous but not defined within a method body (e.g., a synthetic inner class). Every class that has this annotation must also have an `InnerClass` annotation. Additionally, a class may not have both an `EnclosingClass` and an `EnclosingMethod` annotation.

| Name | Format | Description |
|-------|--------|--|
| value | Class | the class which most closely lexically scopes this class |

`dalvik.annotation.EnclosingMethod`

appears on classes

An `EnclosingMethod` annotation is attached to each class which is defined inside a method body. Every class that has this annotation must also have an `InnerClass` annotation. Additionally, a class may not have both an `EnclosingClass` and an `EnclosingMethod` annotation.

| Name | Format | Description |
|-------|--------|---|
| value | Method | the method which most closely lexically scopes this class |

`dalvik.annotation.InnerClass`

appears on classes

An `InnerClass` annotation is attached to each class which is defined in the lexical scope of another class's definition. Any class which has this annotation must also have *either* an `EnclosingClass` annotation *or* an `EnclosingMethod` annotation.

| Name | Format | Description |
|-------------|--------|--|
| name | String | the originally declared simple name of this class (not including any package prefix). If this class is anonymous, then the name is <code>null</code> . |
| accessFlags | int | the originally declared access flags of the class (which may differ from the effective flags because of a mismatch between the execution models of the source language and target virtual machine) |

`dalvik.annotation.MemberClasses`

appears on classes

A `MemberClasses` annotation is attached to each class which declares member classes. (A member class is a direct inner class that has a name.)

| Name | Format | Description |
|-------|---------|-----------------------------|
| value | Class[] | array of the member classes |

`dalvik.annotation.Signature`

appears on classes, fields, and methods

A `Signature` annotation is attached to each class, field, or method which is defined in terms of a more complicated type than is representable by a `type_id_item`. The `.dex` format does not define the format for signatures; it is merely meant to be able to represent whatever signatures a source language requires for successful implementation of that language's semantics. As such, signatures are not generally parsed (or verified) by virtual machine implementations. The signatures simply get handed off to higher-level APIs and tools (such as debuggers). Any use of a signature, therefore, should be written so as not to make any assumptions about only receiving valid signatures, explicitly guarding itself against the possibility of coming across a syntactically invalid signature.

Because signature strings tend to have a lot of duplicated content, a `Signature` annotation is defined as an *array* of strings, where duplicated elements naturally refer to the same underlying data, and the signature is taken to be the concatenation of all the strings in the array. There are no rules about how to pull apart a signature into separate strings; that is entirely up to the tools that generate `.dex` files.

| Name | Format | Description |
|-------|----------|---|
| value | String[] | the signature of this class or member, as an array of strings that is to be concatenated together |

`dalvik.annotation.Throws`

appears on methods

A `Throws` annotation is attached to each method which is declared to throw one or more exception types.

| Name | Format | Description |
|-------|---------|-------------------------------------|
| value | Class[] | the array of exception types thrown |